

A Unified Analysis of Dynamic Interactive Learning

Allerton 2023

Xing Gao, Thomas Maranzatto, Lev Reyzin

University of Illinois Chicago

September 27, 2023

Some Motivation: Recommender Systems

- Recommender systems are integral to many online services today (Amazon, Google, Spotify,...)
- These systems are not static, users preferences change due to external factors
- The system itself may also induce changes in the user
- What is a good mathematical model for this scenario, and can we obtain convergence results?

Previous Work on Static Models

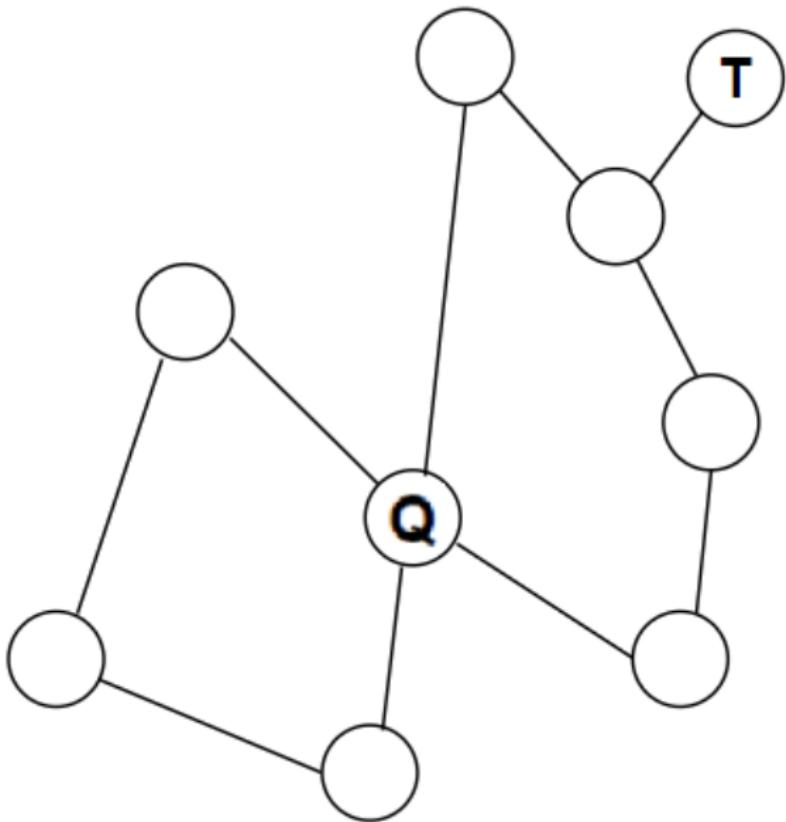
Emamjomeh-Zadeh and Kempe introduced a (static) model for learning with user feedback [1]. It is assumed the user has a finite set of possible preferences modelled as a graph, and a single target preference ranked as ‘the best’.

The learner queries vertices in $v \in V(G)$, and receives feedback from the user as an edge $(v, u) \in E(G)$ on the shortest path from v to t

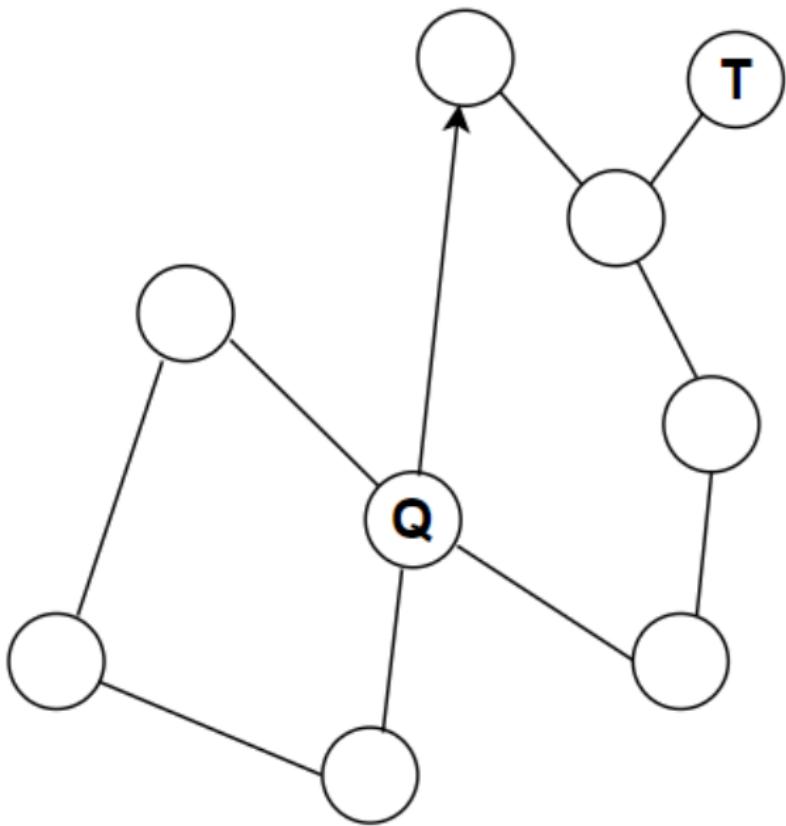
A specification can be given as:

- The learner has access to a weighted graph $G = (V, E, w)$, nodes represent concepts and edges are possible transitions between them
- A target concept $t \in V(G)$ is a fixed node that the learner wishes to discover
- The learner proposes a vertex v as the target, and receives feedback as an edge out of v
- Edge weights satisfy a key property: if the learner proposes $q \neq t$, correct user feedback z lies on a shortest path from q to t
- user feedback may be noisy (and adversarial) with probability $p < \frac{1}{2}$

A Simple Example



A Simple Example



Quick Recap on Weighted Majority

parameter: $\epsilon \in [0, 1]$

Initialize the weights $w_i = 1$ for all experts.

For each round t :

 Make predictions using weighted majority vote based on w .

 For each expert i :

 If the i -th expert's prediction is correct, w_i stays the same.

 Otherwise, $w_i \leftarrow w_i(1 - \epsilon)$.

- Setting: N experts, each makes predictions that are correct some of the time
- Updates bias towards experts which have a history of being correct
- Theorem [2] : MW makes at most $\frac{\log N + m_T^* \log \frac{1}{\epsilon}}{\log \frac{2}{1+\epsilon}}$ mistakes after T rounds, where m_T^* is the number of mistakes made by the best expert

Emamjomeh-Zadeh and Kempe [1] gave a multiplicative weight update algorithm for the digraph problem:

Noiseless case

- Assign initial likelihoods $\lambda(v) = 1$ for each node $v \in G$.
- Query the ‘median’ node with shortest weighted distance to all other nodes: $q = \operatorname{argmin}_v \sum_{v'} d(v, v') \cdot \lambda(v')$.
- Feedback z is a neighbor of the median. $\forall v \in G$:
if v is *consistent* with z , $\lambda(v) \times 1$
if v is *inconsistent* with z , $\lambda(v) \times 0$
 v is *consistent* with z means z is on a shortest $q - v$ path, ie.,
 v could be the target
- $O(\log n)$ queries suffice [1]

A Multiplicative Weights algorithm for Graph Feedback, Noisy Setting

Noisy case

- Feedback is incorrect with probability $p < 1/2$ known to the learner
- Similar weight update idea: after querying median q and receiving feedback z , $\forall v \in G$:
if v is *consistent* with z , $\lambda(v) \times (1 - p)$
if v is *inconsistent* with z , $\lambda(v) \times p$
- with probability $1 - \delta$, target t can be found using

$$\frac{(1 - \delta)}{1 - H(p)} \cdot (\log n + o(\log n) + O(\log^2(1/\delta)))$$

queries. [1]

In a follow-up paper, Emamjomeh-Zadeh and Kempe [3] proposed an extension to the previous paper, where the target is allowed to move B times in the digraph.

- **Shifting-target** model: the target moves in a pre-determined subset of k nodes. The subset is unknown, but the parameter k is known.

Eg: multiple users

- **Drifting-target** model: there's a transition graph G' with maximum degree Δ , and the target moves following the transition graph (G' is known to the learner, and not necessarily the same as the feedback graph G)

Eg: user preference changes over time

Sketch of the algorithm

- Consider all V^R possible sequences of targets throughout the R rounds, and treat these as Experts in the Multiplicative Weights Framework
- Query the ‘median node’ at each round, where this depends on the weight of sequences.
- On feedback, update the weight of *sequences* of nodes, as well as individual nodes.
- This gives polynomial mistake bounds, but at the cost of exponential computation. In the special cases of Drifting and Shifting, [3] give (quasi) polynomial-time algorithms.
- The key observation for speed-up is that we don’t need to keep track of sequence weights, just node weights

Previous Bounds [3]

Shifting Target

$$\frac{1}{1 - H(p)} \cdot \left(k \cdot \log n + B \cdot \log k \right) \\ \leq M(A) \leq \\ \frac{1}{1 - H(p)} \cdot \left(k \log n + (B + 1) \log k + R \cdot H(B/R) \right)$$

Drifting Target

$$\frac{1}{1 - H(p)} \cdot \left(\log n + B \cdot \log \Delta \right) \\ \leq M(A) \leq \\ \frac{1}{1 - H(p)} \cdot \left(\log n + B \cdot \log \Delta + R \cdot H(B/R) \right)$$

Our Contributions

- ① We propose a generalized transition model that includes the Shifting and Drifting models as special cases
- ② We improve the lower bound for the Drifting target model to match the upper bound given in [3]
- ③ We analyze algorithms for low diameter graphs that require $O(1)$ computation per round

A General Model

- Our model is inspired by the Drifting target model over a larger graph G'
- Feedback graph $G = (V, E, w)$
Transition graph $G' = (V', E', \pi)$, in general G' is a bigger graph
- Allow multiple vertices in G' to represent each node in G

Sketch of the algorithm

- Apply a multiplicative weights style algorithm, with ‘experts’ given as sequences of vertices in the transition graph:
 $(V(G))^R$
- Transition graph G' bounds the number of possible sequences there are as $n' \cdot \Delta'^B \cdot \binom{R}{B}$
- Likelihoods in the r ’th round are aggregated from duplicate vertices in the transition graph
- Compute median q_r using likelihoods in feedback graph
- update weights in G' based on feedback and transition information, then aggregate into G .

Let n' , Δ' be the number of vertices and maximum degree of the feedback graph G' .

Theorem

[GMR '23] *The Likelihood Update Algorithm runs in time $O(\Delta' \cdot n' + \text{poly}(n))$, uses space $O(n')$, and makes no more than*

$$\frac{1}{1 - H(p)} \cdot (\log n' + B \cdot \log \Delta' + R \cdot H(B/R))$$

mistakes in expectation.

Simplified Shifting Target Analysis

Corollary

([3], [GMR '23]) In the Shifting Target model, The Likelihood Update Algorithm runs in time $O(k^2 \cdot n^k)$, uses space $O(k \cdot n^k)$, and makes at most

$$\frac{1}{1 - H(p)} \cdot \left(k \cdot \log n + (B + 1) \cdot \log k + R \cdot H(B/R) \right)$$

mistakes in expectation.

Proof.

The transition graph G' consists of $\binom{n}{k}$ disconnected sub-graphs, where each sub-graph is a clique of size k , corresponding to a subset of k vertices in V . Apply the Theorem from the previous slide. □

Theorem

For every n and Δ' , there exists a Drifting Target problem such that every algorithm makes at least

$$\frac{1}{1 - H(1 - p)} \cdot \left(\log n + B \cdot \log \Delta' + R \cdot H(B/R) \right) \\ - o(\log n + B \cdot \log \Delta' + R \cdot H(B/R))$$

mistakes in expectation.

Consider the following simple algorithm:

Algorithm 1 'Follow the Feedback' Procedure for Interactive Learning

$q_1 \leftarrow \operatorname{argmin}_{i \in V} \sum_{j \in V} w(i, j)$ {Start with a 'center' vertex}

for $1 \leq r \leq R$ **do**

$z_r \leftarrow$ feedback from adversary after querying q_r

$q_{r+1} \leftarrow z_r$ {Follow the feedback for next round}

end for

A natural question is, for what classes of graphs and under what transition models does this give polynomial mistake bounds?

A Diameter Bound

Theorem

Suppose G' is as before, and let d be the diameter of G' . Then the 'Follow the feedback' algorithm makes at most

$$\frac{1}{1-p} \cdot \left(dB - \frac{pB}{1-2p} + pR \right)$$

mistakes in expectation.

Proof Idea

- ① Reduce G' to P_d , a path on d vertices
- ② Algorithm 1 makes a mistake anytime the feedback is not the target vertex
- ③ Apply Markov chain hitting time bounds on P_d